



BANK LOAN APPROVAL PREDICTION – LOGISTIC REGRESSION PROJECT

Predict whether a bank should approve a customer's loan
Goal: Predict Approved = 1 or Not Approved = 0 using customer features.

Project Covers:

- Clean dataset loading
- Feature selection
- Train/test split
- Scaling
- Training (Logistic Regression)
- Evaluation (Accuracy, Precision, Recall, F1)
- Confusion Matrix + Classification Report
- Decision Thresholds
- AUC-ROC Curve
- Prediction function

IMPORT LIBRARIES

```
In [4]: # pandas: Used for data handling and analysis
import pandas as pd

# numpy: Used for numerical operations
import numpy as np

# matplotlib.pyplot: Used for basic data visualization
import matplotlib.pyplot as plt

# seaborn: Built on matplotlib, used for advanced & attractive plots
import seaborn as sns

# train_test_split: Used to divide data into training & testing sets
# Ensures model is tested on unseen data
from sklearn.model_selection import train_test_split

# StandardScaler: Used to scale features (mean = 0, std = 1)
# Important for models like Logistic Regression
from sklearn.preprocessing import StandardScaler

# LogisticRegression: A classification algorithm
```

```

# Used to predict binary outcomes (Yes/No, 0/1)
from sklearn.linear_model import LogisticRegression

from sklearn.metrics import (
    accuracy_score, # Accuracy: Out of all predictions, how many were correct
    precision_score, # Precision: Out of all YES predictions, how many were actual YES
    recall_score, # Recall: Out of all actual YES cases, how many did the model predict YES
    f1_score, # F1 Score: A single score that balances Precision and Recall
    confusion_matrix, # A table that shows Correct and wrong predictions (YES vs NO)
    classification_report, # Shows Accuracy, Precision, Recall, and F1 together
    roc_curve, # A graph that shows how well the model separates YES and NO
    roc_auc_score # A number that tells how good the model is at separating classes
);

```

LOAD DATASET

```

In [6]: df = pd.read_csv("bank_loan.csv")

print("Dataset Loaded Successfully!")

print(df.head())

```

Dataset Loaded Successfully!

	Age	Income	CreditScore	EmploymentStatus	Approved
0	59	40358	812	Unemployed	0
1	49	23267	595	Unemployed	0
2	35	102745	619	Salaried	1
3	63	109588	871	Self-employed	0
4	28	58513	648	Salaried	1

```

In [7]: print("===== DATA OVERVIEW =====")

# Shape of the dataset
print("\nShape:", df.shape)

# Data types and info
print("\nDataset Info:")
df.info()

# Check missing values
print("\nMissing Values in Each Column:")
print(df.isnull().sum());

```

===== DATA OVERVIEW =====

Shape: (500, 5)

Dataset Info:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 5 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Age              500 non-null    int64  
 1   Income            500 non-null    int64  
 2   CreditScore       500 non-null    int64  
 3   EmploymentStatus 500 non-null    object  
 4   Approved          500 non-null    int64  
dtypes: int64(4), object(1)
memory usage: 19.7+ KB
```

Missing Values in Each Column:

```
Age             0
Income          0
CreditScore     0
EmploymentStatus 0
Approved        0
dtype: int64
```

```
In [8]: # Count each Employment Status category
print("\nEmployment Status Distribution:")
print(df["EmploymentStatus"].value_counts())

# Preview the last 10 rows
print("\nLast 10 Rows of the Dataset:")
print(df.tail(10));
```

Employment Status Distribution:

```
EmploymentStatus
Salaried      289
Self-employed 136
Unemployed    75
Name: count, dtype: int64
```

Last 10 Rows of the Dataset:

	Age	Income	CreditScore	EmploymentStatus	Approved
490	43	32219	870	Self-employed	0
491	37	20235	662	Self-employed	0
492	46	62929	856	Self-employed	0
493	28	147309	644	Salaried	1
494	49	87444	469	Salaried	1
495	46	159639	655	Salaried	1
496	30	121834	490	Salaried	0
497	46	104555	721	Salaried	1
498	54	93698	306	Unemployed	0
499	61	144450	432	Salaried	0

```
In [9]: # Summary statistics for numerical features
print("\nSummary Statistics:")
print(df.describe())

# Unique values in Employment Status
print("\nUnique Employment Status Values:")
print(df["EmploymentStatus"].unique())
```

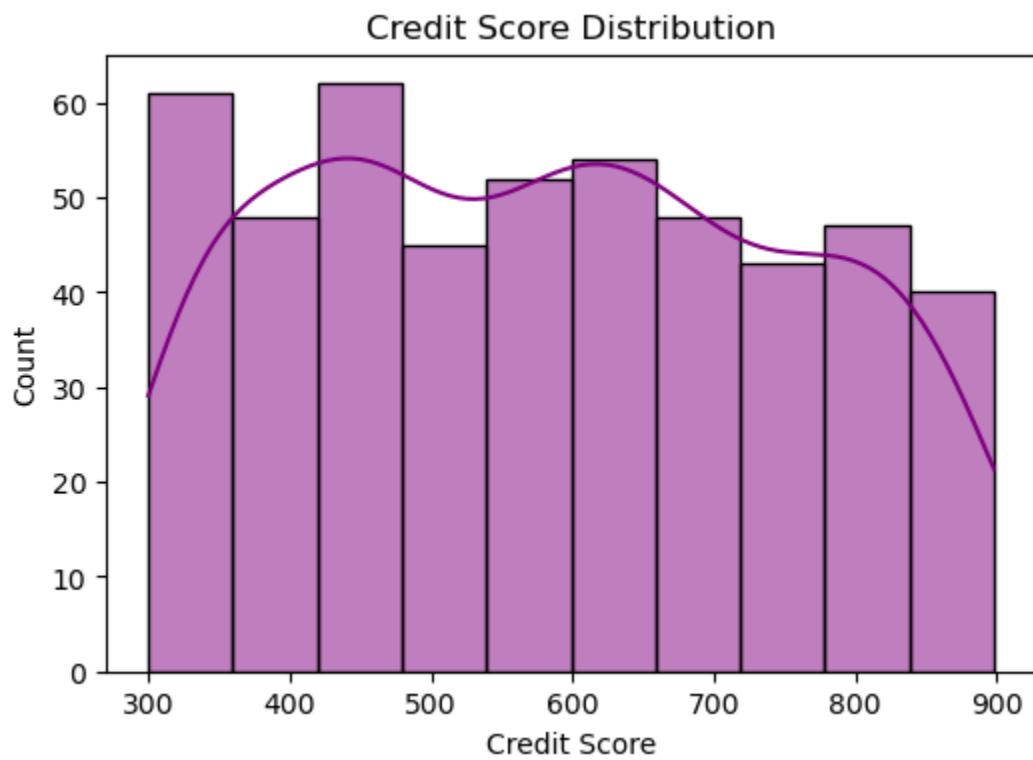
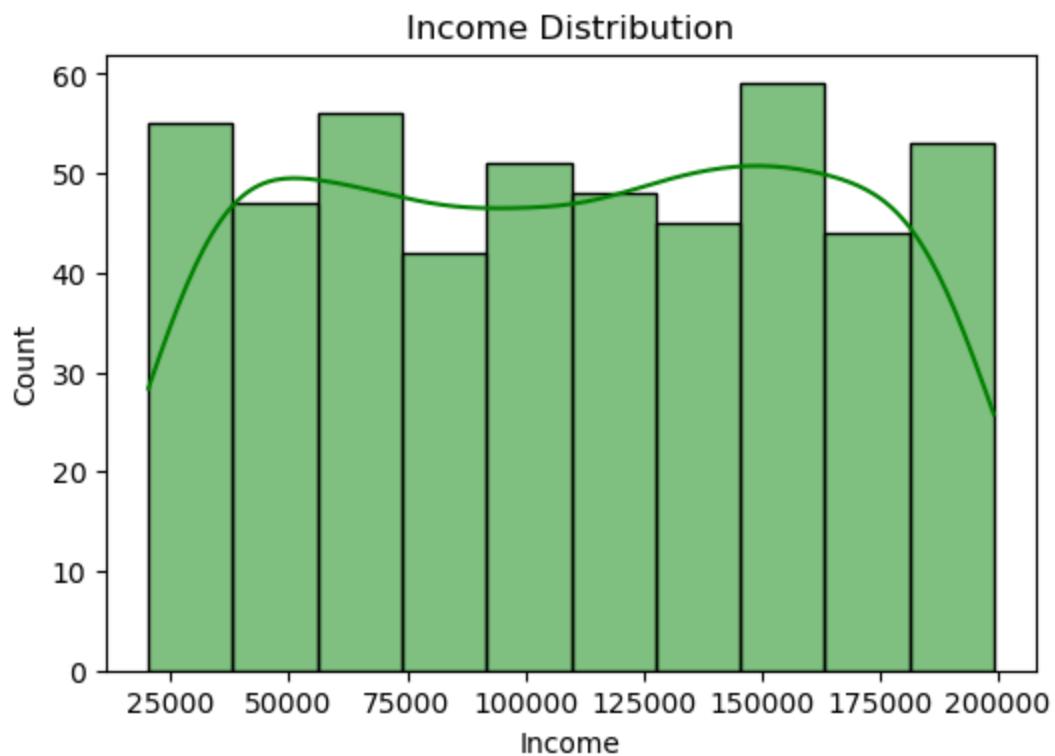
Summary Statistics:

	Age	Income	CreditScore	Approved
count	500.000000	500.000000	500.000000	500.000000
mean	43.116000	109045.696000	583.490000	0.502000
std	12.733217	52974.205023	171.614805	0.500497
min	21.000000	20235.000000	300.000000	0.000000
25%	32.000000	62479.250000	440.000000	0.000000
50%	44.000000	109228.500000	585.500000	1.000000
75%	53.000000	156195.000000	723.500000	1.000000
max	64.000000	199208.000000	898.000000	1.000000

Unique Employment Status Values:

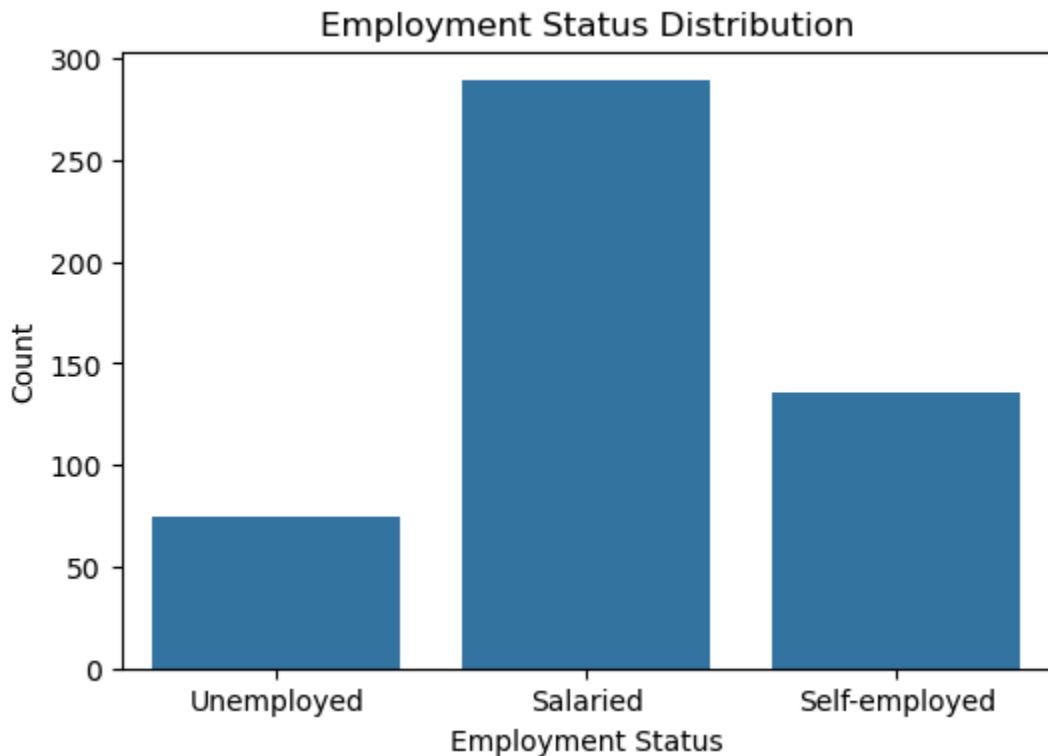
['Unemployed' 'Salaried' 'Self-employed']

```
In [10]: # -----
plt.figure(figsize=(6,4))
sns.histplot(df["Income"], kde=True, color="green")
plt.title("Income Distribution")
plt.xlabel("Income")
plt.ylabel("Count")
plt.show()
# -----
plt.figure(figsize=(6,4))
sns.histplot(df["CreditScore"], kde=True, color="purple")
plt.title("Credit Score Distribution")
plt.xlabel("Credit Score")
plt.ylabel("Count")
plt.show()
```



```
In [11]: # 4. Employment Status Count Plot
plt.figure(figsize=(6,4))
sns.countplot(x="EmploymentStatus", data=df)
plt.title("Employment Status Distribution")
plt.xlabel("Employment Status")
plt.ylabel("Count")
```

```
plt.show()
```



FEATURE SELECTION & CLEANING

```
In [13]: # Select important features  
df = df[["Age", "Income", "CreditScore", "EmploymentStatus", "Approved"]]
```

Encode Employment Status (numeric encoding)

```
In [15]: mapping = {"Unemployed": 0, "Self-employed": 1, "Salaried": 2}  
df["EmploymentEncoded"] = df["EmploymentStatus"].map(mapping);
```

```
In [16]: # Final Feature / Target  
X = df[["Age", "Income", "CreditScore", "EmploymentEncoded"]]  
y = df["Approved"];  
  
print("\nFinal Features:\n", X.head())
```

Final Features:

	Age	Income	CreditScore	EmploymentEncoded
0	59	40358	812	0
1	49	23267	595	0
2	35	102745	619	2
3	63	109588	871	1
4	28	58513	648	2

TRAIN-TEST SPLIT

```
In [18]: X_train, X_test, y_train, y_test = train_test_split(  
    X, y, test_size=0.25, random_state=42, stratify=y  
);  
print("Train size:", X_train.shape)  
print("Test size:", X_test.shape)
```

Train size: (375, 4)
Test size: (125, 4)

FEATURE SCALING

```
In [20]: scaler = StandardScaler()  
# Create a scaler that standardizes data (mean = 0, std = 1)  
  
X_train_scaled = scaler.fit_transform(X_train)  
# Learn scaling from training data and apply it  
  
X_test_scaled = scaler.transform(X_test)  
# Apply the SAME scaling to test data (no learning)
```

TRAIN LOGISTIC REGRESSION MODEL

```
In [22]: model = LogisticRegression(max_iter=1000) # Improves accuracy  
# Create a Logistic Regression model  
  
model.fit(X_train_scaled, y_train)  
# Train the model using the scaled training data and their correct labels  
  
print("Model Training Completed!")  
# Confirm that training is finished
```

Model Training Completed!

MODEL PREDICTIONS

```
In [24]: y_pred = model.predict(X_test_scaled)  
# Predict class labels (0 or 1) for the test data  
  
y_prob = model.predict_proba(X_test_scaled)[:, 1]  
# Get probability of class 1 (loan approved) for each test sample
```

EVALUATION METRICS

```
In [69]: print("\n===== MODEL EVALUATION =====")  
  
print("Accuracy :", accuracy_score(y_test, y_pred));  
# How many total predictions were correct
```

```

print("Precision:", precision_score(y_test, y_pred));
# Out of all predicted YES, how many were actually YES (controls false positive)

print("Recall    :", recall_score(y_test, y_pred));
# Out of all actual YES, how many the model correctly found (controls false negative)

print("F1 Score :", f1_score(y_test, y_pred));
# Balance of Precision and Recall in one score

```

===== MODEL EVALUATION =====

```

Accuracy : 0.624
Precision: 0.6212121212121212
Recall   : 0.6507936507936508
F1 Score : 0.6356589147286822

```

Model Evaluation

Accuracy = 0.62 (62%)

Out of 100 loan predictions, about 62 are correct.

Precision = 0.62 (62%)

When the model says "YES – loan approved", it is correct 62 times out of 100.

Recall = 0.65 (65%)

Out of all people who should get the loan, the model correctly finds 65% of them.

F1 Score = 0.63 (63%)

Overall balanced performance between Precision and Recall.

Percentages help us understand model performance easily, regardless of dataset size.

CONFUSION MATRIX + CLASSIFICATION REPORT

```
In [78]: print("\n===== CONFUSION MATRIX =====")
print(confusion_matrix(y_test, y_pred));
# Shows TP, FP, TN, FN to understand correct and incorrect predictions
```

===== CONFUSION MATRIX =====

```
[[37 25]
 [22 41]]
```

Confusion Matrix

37 → Correctly predicted NO

25 → Wrongly predicted YES (but actually NO)
22 → Wrongly predicted NO (but actually YES)
41 → Correctly predicted YES

```
In [83]: print("\n===== CLASSIFICATION REPORT =====")  
print(classification_report(y_test, y_pred));  
# Shows Precision, Recall, F1-score, and support for each class
```

===== CLASSIFICATION REPORT =====				
	precision	recall	f1-score	support
0	0.63	0.60	0.61	62
1	0.62	0.65	0.64	63
accuracy			0.62	125
macro avg	0.62	0.62	0.62	125
weighted avg	0.62	0.62	0.62	125

Class 0 (NO / Loan Rejected)

Precision 63% → When model says NO, it is correct 63 times out of 100
Recall 60% → Found 60% of actual NO cases
F1-score 61% → Balanced score for NO class
Support 62 → Total 62 NO records

Class 1 (YES / Loan Approved)

Precision 62% → When model says YES, it is correct 62 times out of 100
Recall 65% → Found 65% of actual YES cases
F1-score 64% → Balanced score for YES class
Support 63 → Total 63 YES records

Accuracy = 62% → Model predicted 62 correct out of 100
Macro Avg → Simple average of both classes
Weighted Avg → Average based on number of records

Decision Threshold Analysis

Decision Threshold Analysis means deciding at what probability value the model should say YES or NO.

≥ 0.5 → YES
< 0.5 → NO

```
In [88]: threshold = 0.5; # Equal importance to YES and NO | Balanced datasets
print("\nDefault Threshold =", threshold)
print("Prediction Probabilities (first 10):")
print(y_prob[:20]); # Shows the first 10 probability values predicted by the
```

Default Threshold = 0.5
Prediction Probabilities (first 10):
[0.32716288 0.68291553 0.52806798 0.675654 0.38418969 0.71675992
 0.69488846 0.63420475 0.59271349 0.5549581 0.45405383 0.57621246
 0.718221 0.08868347 0.85116035 0.36293882 0.27502011 0.64420197
 0.18260534 0.24739485]

AUC-ROC Curve

AUC-ROC tells us how well a model can separate positive and negative cases.

It measures how good the model is at distinguishing YES vs NO.

```
In [95]: # AUC tells how well the model separates Yes vs No.

# Step 1: Calculate False Positive Rate, True Positive Rate and Thresholds
fpr, tpr, thresholds = roc_curve(y_test, y_prob)
# fpr = how many wrong positives at each threshold
# tpr = how many correct positives at each threshold
# thresholds = different probability cut-off values

# Step 2: Calculate AUC score
auc = roc_auc_score(y_test, y_prob)
# AUC tells how well the model separates class 0 and class 1
# Higher AUC = better model

print("\nAUC Score:", auc)
```

AUC Score: 0.7073732718894009

AUC Score = 0.71

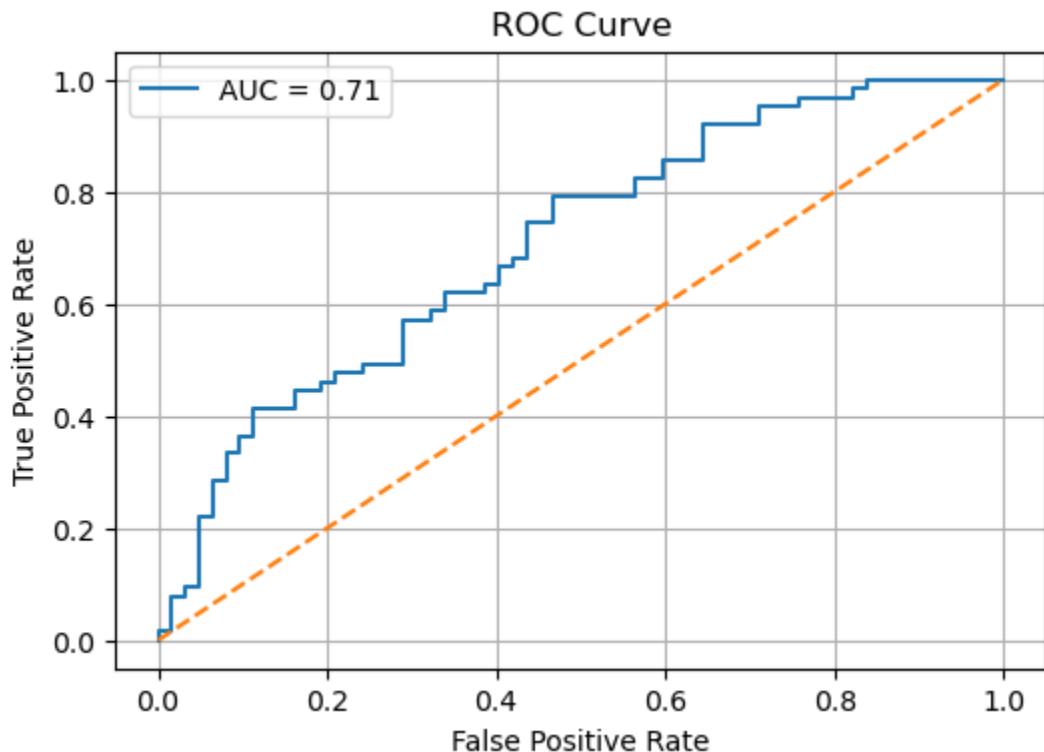
Means 71% chance the model will correctly separate YES vs NO

```
In [99]: plt.figure(figsize=(6,4))
plt.plot(fpr, tpr, label=f"AUC = {auc:.2f}")
# Plot the ROC curve

plt.plot([0,1], [0,1], '--') # Draws a random guessing line | Used for comparison

plt.xlabel("False Positive Rate");
plt.ylabel("True Positive Rate");
```

```
plt.title("ROC Curve")
plt.legend()
plt.grid(True)
plt.show()
```



PREDICTION FUNCTION

```
In [119]: def predict_loan(age, income, credit, employment):

    # 1. Put the customer data into a DataFrame
    row = pd.DataFrame([[age, income, credit, mapping[employment]]],
                       columns=["Age", "Income", "CreditScore", "EmploymentEnc"])

    # 2. Scale the data (same scaling used during training)
    row_scaled = scaler.transform(row);

    # 3. Get probability of loan approval
    prob = model.predict_proba(row_scaled)[0][1]

    # 4. Final prediction (0 or 1) using threshold 0.5
    pred = 1 if prob >= 0.5 else 0;

    # 5. Print simple output
    print(f"\nLoan Approval Probability = {prob:.2f}")

    if pred == 1:
        print("Loan Approved")
    else:
```

```
print("Loan Not Approved")  
  
predict_loan(25, 60000, 800, "Salaried")  
predict_loan(55, 30000, 500, "Unemployed")
```

Loan Approval Probability = 0.57
Loan Approved

Loan Approval Probability = 0.12
Loan Not Approved

In []:

In []: